

CSci 5271: Introduction to Computer Security

Exercise Set 3

due: November 6, 2012

Ground Rules. You may choose to complete these exercises in a group of up to three students. Each group should turn in **one** copy with the names of all group members on it. You may use any source you can find to help with this assignment but you **must** explicitly reference any source you use besides the lecture notes or textbook.

1. TCP-Unfriendly. TCP’s “congestion control” mechanism relies on *users* to respond appropriately to network congestion by backing off their sending rate. One potential problem with this mechanism is what’s called (by economists) the “tragedy of the commons:” suppose Alice knows that everyone else obeys TCP’s congestion control mechanism. Then if she continues sending at the same rate, everyone else will slow down a little bit more and she will get better service from the network. So Alice has no motivation to obey TCP congestion control (other than the fact that not doing so involves finding or writing her own TCP stack - *details, details*) and in fact neither does anyone else. But if no one obeys the mechanism, the network (commons) becomes useless, **a tragedy**.

- (a) Bob the Network Builder has an idea about how to solve this problem. He reasons that congested routers can see the *exact* state of a TCP connection. So if a particular connection does not slow down in response to dropped packets, the router can send a **RST** packet to each end of the connection. This will cause both ends of the connection to drop the connection, much more painful than just dropping an odd packet or two. From a *security* standpoint, what’s the problem with Bob’s idea – that is, if I’m an adversary intent on sending at a high rate, can I circumvent this mechanism?
- (b) When Bob realizes that reset packets aren’t sufficient, he proposes a more direct attack: *blacklisting*. Under this idea, routers that notice TCP senders that don’t respond to dropped packets appropriately will just stop routing packets for that sender. List several ways in which this is both ineffective against adversaries and a generally bad idea if adversaries got wind of it.

2. The MAC Daddy Exercise. Armed with a copy of Schneier’s *Applied Cryptography* from a used bookstore, Sly can’t wait to design his own encrypted thingamadoodad protocol. He starts off with an uber-secure key exchange protocol (we’ll talk more about this in the next exercise) that ends with Alice and Bob sharing secret keys for encryption (K_e) and authentication (K_a). Now he wants to design a secure symmetric channel using these keys.

- (a) Sly decides at first that he wants to use AES-CBC-MAC (with 128 bit blocks) for integrity. He looks carefully at his key exchange protocol and realizes that an adversary can interfere to make Alice and Bob end up deciding on different keys. So the first message sent over by Alice will be $\tau_0 = \text{cbcMAC}_{K_a}(0^{128}) = \text{aesEncrypt}_{K_a}(0^{128})$. If Bob’s local value doesn’t check out, he aborts, otherwise the channel is usable. Afterwards, whenever Alice wants to send the message M over the secure channel, she’ll compute $\tau_M \leftarrow \text{cbcMAC}_{K_a}(M)$ and send the pair (M, τ_M) over the channel; Bob will check whether $\tau_M = \text{cbcMAC}_{K_a}(M)$ and if so will conclude that Alice said M .

This is a pretty bad idea. Show how to use the values τ_0 , M and τ_M to compose a message to Bob that will convince him Alice meant to say the two-block message (M, τ_M) instead of just M . Explain why your message will convince Bob that Alice meant to say (M, τ_M) rather than just M .¹ Since τ_M is just 128 random-looking bits, why is this a big deal?

- (b) Sly's friend Sally notices the same attack on his scheme. She proposes a different method of authenticating (and encrypting) messages: ignore the key K_a . Instead, to authenticate and encrypt the message M , first compute $H(M)$ using SHA-256; then encrypt $(M, H(M))$ together, using AES-CTR encryption. So the message sent on the insecure channel would be $\text{CTR} - \text{Encrypt}_{K_e}(M, H(M))$; Bob would decrypt the message using K_e , check that the last 256 bits of the plaintext are the hash of all of the previous bits, and accept the message if they are.

Show that this is also a bad idea: if Alice ever sends a ciphertext corresponding to the message M , where Eve knows M , Eve can generate a ciphertext corresponding to any message M' , (of the same length as M) that Bob will accept. (For example, if Alice sends the message "ATTACK AT TEN AM" Eve can drop it and make Bob accept the message "GO BACK HOME BOB" instead.)

3. Protocol Droid. Artoo and C3P0 often fly on different starships and need to alert each other to their presence when their ships come in contact - otherwise they might accidentally blow each other up! They agree on a shared key K and a MAC algorithm that outputs t bit tags to use in the following protocol.

1. $A \rightarrow C$: a random t -bit string N_A and $\text{MAC}_K(N_A)$.
2. C : on message n, t checks that $\text{MAC}_K(n) = t$, and if so, he accepts A , otherwise he blows up the other party.
3. $C \rightarrow A$: $\text{MAC}_K(t)$.
4. A : on message t' checks that $t' = \text{MAC}_K(\text{MAC}_K(N_A))$. If so, he accepts C , otherwise he blows him up.

The idea being that A proves he is A by correctly MACing N_A (which, if the key is secret, only A or C could do) and C proves he is C by MACing the MAC. But...

- (a) A and C use this protocol for a while and then discover, to their dismay, that sometimes the Emperor, E , has been successfully fooling C into believing he is A . Supposing that middleperson attacks (and "using the force") are prevented by speed-of-light limitations, what is a simple way for him to do this?
- (b) A and C decide that one way to prevent the attack is for C to remember every value of N_A used in a previous challenge and reject if one is ever reused. Suppose E sees one authentication between A and C . How can he fool C into believing he is A as many times afterwards as he wants?

¹Hint: for any string A , $A \oplus A = 0^{|A|}$ and $A \oplus 0^{|A|} = A$. One more hint: try writing a recursive definition of CBC-MAC

4. Hashing and Signing. Nearly every digital signature scheme works by first hashing a message to be signed (with a cryptographic hash function) and then performing some operation on the hash - so in essence, we are “signing the hash” and not the message. In particular, if Eve sees Alice’s signature on the message M and can find a message $M' \neq M$ so that $H(M) = H(M')$, she can convince people that Alice signed M' . This is OK, since a good crypto hash function H will resist targeted collisions.

Suppose our signature scheme uses a hash function H with an output length ℓ that is sufficient to resist targeted collisions but NOT resistant to free collisions (e.g. the hash length is around 100-120 bits). Then it is possible that if Eve can get Alice to sign one of a pair of colliding messages, she can later claim that Alice signed the other.

- (a) It might be tempting to think that the risk of such an attack is minimal, since the birthday attack works by hashing random messages until two have the same hash; why would Alice want to sign a random message, and even if she did, why would we care that Eve could claim she signed a different random message? Give a simple explanation why even this attack could be troublesome.
- (b) Let’s show that in fact, it is worse than that. Suppose that a message is “favorable” if it is something that Alice would sign, for example “I will pay \$5 to McDonald’s for my lunch.” Suppose that a message is “undesirable” if it is something that Alice would not sign, like “I will pay \$500,000 to Eve for her lunch.” Notice that we can generate 256 different “favorable” messages from the example above, just by varying the number of space characters between words (just switching between 1 and 2). Extend this idea to show how to generate a pair of messages, one favorable and one undesirable, with the same hash. Your attack should compute about as many hashes as the birthday attack.
- (c) Complete the attack: how would Eve use the pair she generates in part (b) to her advantage?